

Implementation

1. Structure

The program implement rendering in cuda by converting `render` to a kernel, which renders one pixel per threads. In addition, the `3x3` anti-aliasing loops' first layer is unrolled by specifying `m` to the kernel and launching it 3 times. Then, another kernel `_convert_pixel` is used to convert floating point colors to uint8 pixel values.

2. Parrtition

Each block has 2D $32 \times 8 = 256$ threads, corresponding to a 32×8 pixels region. And the whole image is splitted to $\text{ceil}(w/32) \times \text{ceil}(h/8)$ blocks.

3. Optimization

The main optimization is to turn on `-use_fast_math` flag in the compiler and use float instead of double. Then, I also converted many operations to their fused version. For example:

- `A * B + C -> __fmaf(A, B, C)`
- `s = sin(T), c = cos(T) -> __sincosf(T, &s, &c)`

Another small trick is to turn off PNG compression for reduced sequential runtime.

Analysis

1. GPU kernel execution

Kernel	Avg	Min	Max	Time(%)	Time	Calls
<code>render_pixel</code>				98.98%	2.33286s	3
	777.62ms	763.78ms	791.09ms			
<code>Copy device buffer to host</code>				0.98%	23.175ms	1
	23.175ms	23.175ms	23.175ms			
<code>convert_pixel</code>				0.02%	456.09us	1
	456.09us	456.09us	456.09us			
<code>Initialize float device buffer</code>				0.01%	310.24us	1
	310.24us	310.24us	310.24us			
<code>Copy global varaibles to device</code>				0.00%	8.3520us	9
	928ns	896ns	960ns			

2. Nsight Compute

Under different settings of threads (`32x4`, `32x8`, `32x16`), the throughput of SM all reach ~80%, so the computation is bounded by ALU and adjusting the partition won't help.

3. Additional analysis

IO can help a lot.

Conclusion

1. It's hard to optimize the arithmetic operations.
2. You are amazing.